

Data Mining Project

Gautam

2024-10-25

This Project includes healthcare cost prediction for Medicare patients specially for Heart Failure, where i am using two different model Linear regression and Decision tree to predict the cost for Heart Failure and compare among the model result.

Description of the Code Process

1. Load Libraries and Data
2. Merge Datasets and Identify Missing Values
3. Calculate and Visualize Correlation Matrix
4. Collect Categorical Variables and Perform ANOVA
5. Select Features and Filter by DRG Codes
6. Convert Categorical Variables to Factor and Split Data
7. Preprocess Data and Create Linear Model with Features
8. Make Predictions and Calculate Evaluation Metrics
9. Print Evaluation Results and Lasso Model

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
##
## intersect, setdiff, setequal, union
```

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v forcats 1.0.0 v readr 2.1.5
## v ggplot2 3.4.4 v stringr 1.5.1
## v lubridate 1.9.3 v tibble 3.2.1
## v purrr 1.0.2 v tidyr 1.3.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(ggplot2)
library(maps)
```

```
##
## Attaching package: 'maps'
##
## The following object is masked from 'package:purrr':
##
## map
```

```
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
## lift
```

```
library(rpart)
library(rpart.plot)
library(rattle)
```

```
## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(neuralnet)
```

```
##
## Attaching package: 'neuralnet'
```

```
##
## The following object is masked from 'package:dplyr':
##
##      compute
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 4.3.3
```

```
## corrplot 0.95 loaded
```

```
library(reshape2)
```

```
##
## Attaching package: 'reshape2'
##
## The following object is masked from 'package:tidyr':
##
##      smiths
```

```
# Load the data
```

```
hospital_data <- read.csv("/Users/Lord/Downloads/Hospital_General_Information.csv")
```

```
medicare_data <- read.csv("/Users/Lord/Downloads/Medicare_Inpatient_Hospital_by_Provider_and_Service_20
```

```
# Rename ZIP.Code column as it is ZIP.Code in hospital data
```

```
colnames(medicare_data)[7] <- "ZIP.Code"
```

```
# Merging datasets on ZIP.Code
```

```
merged_data <- inner_join(medicare_data, hospital_data, "ZIP.Code")
```

```
## Warning in inner_join(medicare_data, hospital_data, "ZIP.Code"): Detected an unexpected many-to-many
## i Row 489 of 'x' matches multiple rows in 'y'.
## i Row 1 of 'y' matches multiple rows in 'x'.
## i If a many-to-many relationship is expected, set 'relationship =
##      "many-to-many"' to silence this warning.
```

```
# Identifying missing values
```

```
miss_val <- sum(is.na(merged_data))
```

```
# Calculating correlation matrix for numeric features to identify highly variable and related features
```

```
numeric_data <- merged_data %>% select_if(is.numeric)
```

```
correlation_matrix <- cor(numeric_data, use = "pairwise.complete.obs")
```

```
# Creating a heatmap of correlation matrix
```

```
melted_corr <- melt(correlation_matrix)
```

```
ggplot(data = melted_corr, aes(x = Var1, y = Var2, fill = value)) +
```

```
  geom_tile() +
```

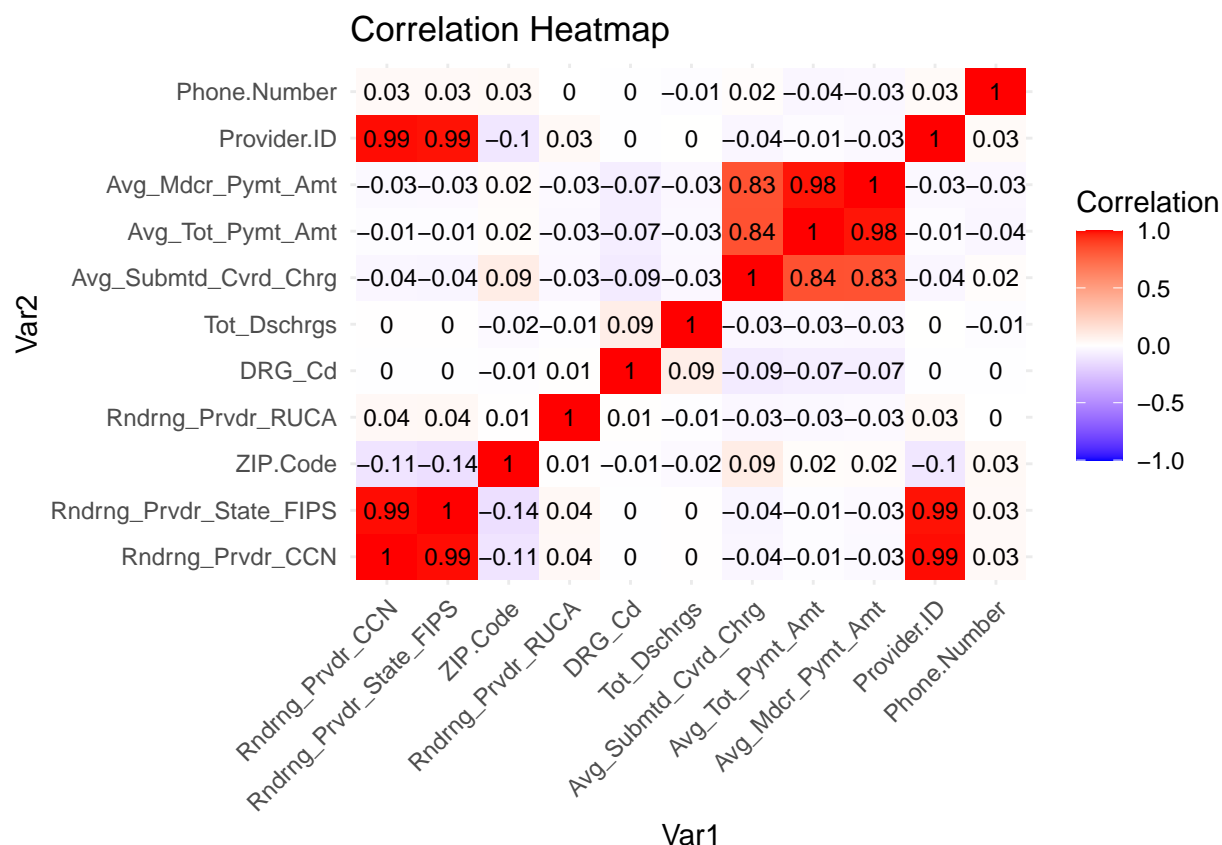
```
  scale_fill_gradient2(low = "blue", high = "red", mid = "white", limit = c(-1, 1), name = "Correlation
```

```
  theme_minimal() +
```

```
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
```

```
  labs(title = "Correlation Heatmap") +
```

```
  geom_text(aes(label = round(value, 2)), color = "black", size = 3)
```



```
#Collecting only categorical variables
categorical_vars <- names(merged_data)[sapply(merged_data, function(x) is.factor(x) || is.character(x))]
print(categorical_vars)
```

```
## [1] "Rndrng_Privr_Org_Name"
## [2] "Rndrng_Privr_St"
## [3] "Rndrng_Privr_City"
## [4] "Rndrng_Privr_State_Abrvtn"
## [5] "Rndrng_Privr_RUCA_Desc"
## [6] "DRG_Desc"
## [7] "Hospital.Name"
## [8] "Address"
## [9] "City"
## [10] "State"
## [11] "County.Name"
## [12] "Hospital.Type"
## [13] "Hospital.Ownership"
## [14] "Emergency.Services"
## [15] "Meets.criteria.for.meaningful.use.of.EHRs"
## [16] "Hospital.overall.rating"
## [17] "Hospital.overall.rating.footnote"
## [18] "Mortality.national.comparison"
## [19] "Mortality.national.comparison.footnote"
## [20] "Safety.of.care.national.comparison"
## [21] "Safety.of.care.national.comparison.footnote"
## [22] "Readmission.national.comparison"
```

```
## [23] "Readmission.national.comparison.footnote"
## [24] "Patient.experience.national.comparison"
## [25] "Patient.experience.national.comparison.footnote"
## [26] "Effectiveness.of.care.national.comparison"
## [27] "Effectiveness.of.care.national.comparison.footnote"
## [28] "Timeliness.of.care.national.comparison"
## [29] "Timeliness.of.care.national.comparison.footnote"
## [30] "Efficient.use.of.medical.imaging.national.comparison"
## [31] "Efficient.use.of.medical.imaging.national.comparison.footnote"
```

```
#Selecting only categorical variables to perform anova test for selecting highly variable feature to pr
anova_result <- aov(Avg_Mdcr_Pymt_Amt ~ Hospital.Type + Hospital.Ownership+ State + Hospital.overall.
summary(anova_result)
```

```
##
##
## Df Sum Sq Mean Sq
## Hospital.Type 2 3.388e+11 1.694e+11
## Hospital.Ownership 9 3.165e+11 3.516e+10
## State 50 1.036e+12 2.072e+10
## Hospital.overall.rating 5 9.291e+10 1.858e+10
## Emergency.Services 1 1.354e+08 1.354e+08
## Meets.criteria.for.meaningful.use.of.EHRs 1 1.323e+09 1.323e+09
## Mortality.national.comparison 3 2.310e+11 7.699e+10
## Patient.experience.national.comparison 3 3.105e+10 1.035e+10
## Timeliness.of.care.national.comparison 3 3.262e+11 1.087e+11
## Efficient.use.of.medical.imaging.national.comparison 3 3.167e+10 1.056e+10
## Residuals 250952 4.858e+13 1.936e+08
##
## F value Pr(>F)
## Hospital.Type 875.061 < 2e-16 ***
## Hospital.Ownership 181.665 < 2e-16 ***
## State 107.044 < 2e-16 ***
## Hospital.overall.rating 95.998 < 2e-16 ***
## Emergency.Services 0.700 0.40289
## Meets.criteria.for.meaningful.use.of.EHRs 6.836 0.00893 **
## Mortality.national.comparison 397.747 < 2e-16 ***
## Patient.experience.national.comparison 53.470 < 2e-16 ***
## Timeliness.of.care.national.comparison 561.810 < 2e-16 ***
## Efficient.use.of.medical.imaging.national.comparison 54.530 < 2e-16 ***
## Residuals
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Selecting Features based on Correlation Matrix and anova result
selected_vars <- c("Avg_Mdcr_Pymt_Amt", "Avg_Tot_Pymt_Amt", "Hospital.Type", "Hospital.Ownership",
                  "Hospital.overall.rating", "Meets.criteria.for.meaningful.use.of.EHRs",
                  "Mortality.national.comparison", "Patient.experience.national.comparison",
                  "Timeliness.of.care.national.comparison", "Efficient.use.of.medical.imaging.national
```

```
#Filtering dataset for HeartFailure DRG_Cd Code
final_data_HF <- filter(merged_data, DRG_Cd %in% c("291", "292", "293"))
selected_data_HF <- final_data_HF %>% select(all_of(selected_vars))
```

```
#Converting categorical variables to factor to pass in model- Used as optional for one-hot encoding
```

```

variables_to_factor <- c("Hospital.Type", "Hospital.Ownership", "Hospital.overall.rating",
                        "Meets.criteria.for.meaningful.use.of.EHRs", "Mortality.national.comparison",
                        "Patient.experience.national.comparison",
                        "Timeliness.of.care.national.comparison",
                        "Efficient.use.of.medical.imaging.national.comparison")

selected_data_HF[variables_to_factor] <- lapply(selected_data_HF[variables_to_factor], as.factor)

# Train-test split
set.seed(123)
train_index <- createDataPartition(selected_data_HF$Avg_Mdcr_Pymt_Amt, p = 0.8, list = FALSE)
train_data <- selected_data_HF[train_index, ]
test_data <- selected_data_HF[-train_index, ]

# Preprocessing data
pre_process <- preProcess(train_data %>% select(-Avg_Mdcr_Pymt_Amt), method = c("center", "scale"))
train_data_normalized <- predict(pre_process, newdata = train_data %>% select(-Avg_Mdcr_Pymt_Amt))
test_data_normalized <- predict(pre_process, newdata = test_data %>% select(-Avg_Mdcr_Pymt_Amt))

# Adding the target variable back
train_data_normalized$Avg_Mdcr_Pymt_Amt <- train_data$Avg_Mdcr_Pymt_Amt
test_data_normalized$Avg_Mdcr_Pymt_Amt <- test_data$Avg_Mdcr_Pymt_Amt

# Creating linear model with interaction and polynomial features for capturing non-linearity and better
linear_model_cv <- train(Avg_Mdcr_Pymt_Amt ~ . + I(Avg_Tot_Pymt_Amt^2) +
                        Avg_Tot_Pymt_Amt:Hospital.Type +
                        Avg_Tot_Pymt_Amt:Hospital.Ownership,
                        data = train_data_normalized, method = "lm",
                        trControl = trainControl(method = "repeatedcv", number = 10, repeats = 3))

## Warning in predict.lm(modelFit, newdata): prediction from rank-deficient fit;
## attr(*, "non-estim") has doubtful cases

## Warning in predict.lm(modelFit, newdata): prediction from rank-deficient fit;
## attr(*, "non-estim") has doubtful cases

## Warning in predict.lm(modelFit, newdata): prediction from rank-deficient fit;
## attr(*, "non-estim") has doubtful cases

## Warning in predict.lm(modelFit, newdata): prediction from rank-deficient fit;
## attr(*, "non-estim") has doubtful cases

predictions <- predict(linear_model_cv, newdata = test_data_normalized)

# Calculating evaluation metrics
train_predictions <- predict(linear_model_cv, newdata = train_data_normalized)
train_rmse <- RMSE(train_predictions, train_data_normalized$Avg_Mdcr_Pymt_Amt)
test_rmse <- RMSE(predictions, test_data_normalized$Avg_Mdcr_Pymt_Amt)
train_R2 <- R2(train_predictions, train_data_normalized$Avg_Mdcr_Pymt_Amt)

```

```
test_R2 <- R2(predictions, test_data_normalized$Avg_Mdcr_Pymt_Amt)
```

```
# Creating Dataframe of the result
evaluation_results <- data.frame(
  RMSE = c(train_rmse, test_rmse),
  R2 = c(train_R2, test_R2),
  row.names = c("Train", "Test")
)
print(evaluation_results)
```

```
##           RMSE           R2
## Train 800.2127 0.9448077
## Test  718.0019 0.9465313
```

```
# Creating a scatter plot of actual vs. predicted values
results_df <- data.frame(
  Actual = test_data$Avg_Mdcr_Pymt_Amt,
  Predicted = predictions
)

ggplot(results_df, aes(x = Actual, y = Predicted)) +
  geom_point(color = "blue", alpha = 0.5) +
  geom_abline(slope = 1, intercept = 0, color = "red") +
  labs(title = "Actual vs Predicted Values",
       x = "Actual Values",
       y = "Predicted Values") +
  theme_minimal() +
  xlim(0, max(results_df$Actual) * 1.1) +
  ylim(0, max(results_df$Predicted) * 1.1)
```

Actual vs Predicted Values



```
#Lasso Model to check overfitting or possible mutiple colinearity
lasso_model <- train(Avg_Mdcr_Pymt_Amt ~ ., data = train_data_normalized,
                      method = "glmnet",
                      trControl = trainControl(method = "cv", number = 10),
                      tuneGrid = expand.grid(alpha = 1, lambda = seq(0.01, 1, length.out = 100)),
                      print = FALSE)
print(lasso_model)
```

```
## glmnet
##
## 6171 samples
##    9 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 5553, 5552, 5555, 5554, 5554, 5554, ...
## Resampling results across tuning parameters:
##
##   lambda  RMSE      Rsquared  MAE
##   0.01    809.231  0.9423251  494.8502
##   0.12    809.231  0.9423251  494.8502
##   0.23    809.231  0.9423251  494.8502
##   0.34    809.231  0.9423251  494.8502
##   0.45    809.231  0.9423251  494.8502
##   0.56    809.231  0.9423251  494.8502
##   0.67    809.231  0.9423251  494.8502
##   0.78    809.231  0.9423251  494.8502
```



```
##    0.89    809.231  0.9423251  494.8502
##    1.00    809.231  0.9423251  494.8502
##
## Tuning parameter 'alpha' was held constant at a value of 1
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 1 and lambda = 1.
```

#Decision Tree

Training Decision Tree Model with Cross-Validation

```
tree_model <- train(
  Avg_Mdcr_Pymt_Amt ~ . + I(Avg_Tot_Pymt_Amt^2) + Avg_Tot_Pymt_Amt:Hospital.Type + Avg_Tot_Pymt_Amt:Hosp
  data = train_data_normalized,
  method = "rpart",
  trControl = trainControl(method = "cv", number = 10),
  tuneGrid = expand.grid(cp = seq(0.001, 0.05, by = 0.005))
)
```

```
predictions_tree <- predict(tree_model, newdata = test_data_normalized)
train_predictions_tree <- predict(tree_model, newdata = train_data_normalized)
```

Calculating evaluation metrics

```
rmse_train_tree <- RMSE(train_predictions_tree, train_data_normalized$Avg_Mdcr_Pymt_Amt)
r2_train_tree <- R2(train_predictions_tree, train_data_normalized$Avg_Mdcr_Pymt_Amt)
```

```
rmse_test_tree <- RMSE(predictions_tree, test_data_normalized$Avg_Mdcr_Pymt_Amt)
r2_test_tree <- R2(predictions_tree, test_data_normalized$Avg_Mdcr_Pymt_Amt)
```

#creating dataframe of the error

```
tree_error_df <- data.frame(
  Metric = c("RMSE", "R-squared"),
  Training_Error = c(rmse_train_tree, r2_train_tree),
  Testing_Error = c(rmse_test_tree, r2_test_tree)
)
```

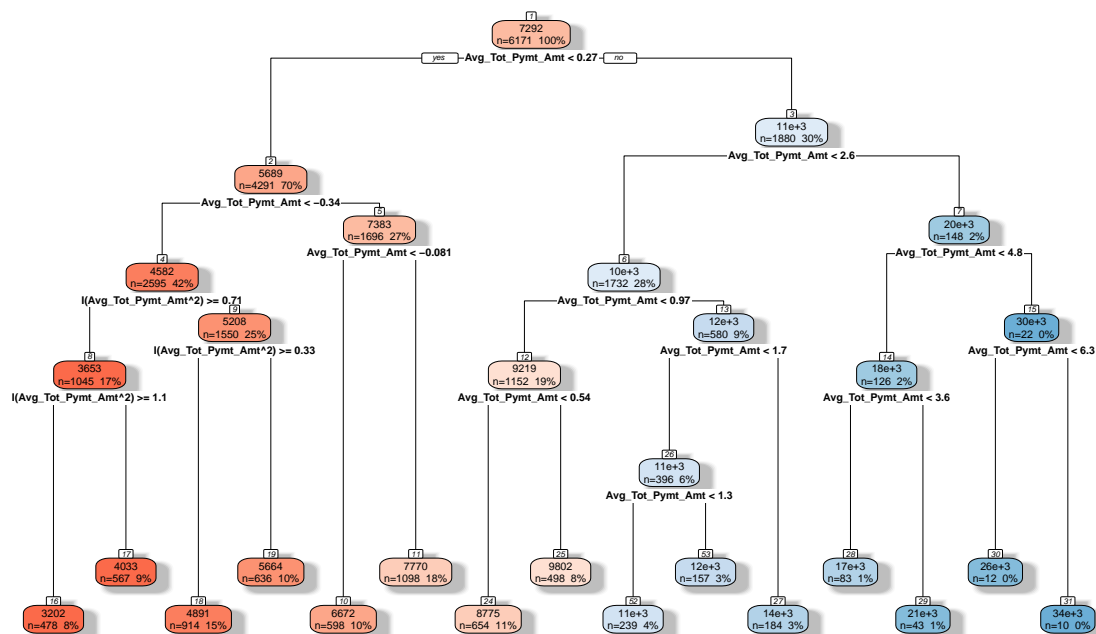
```
print(tree_error_df)
```

```
##      Metric Training_Error Testing_Error
## 1      RMSE      893.1954395      793.6572309
## 2 R-squared       0.9312361       0.9344725
```

Plotting the decision tree for regression

```
rpart.plot(tree_model$finalModel, type = 2, extra = 101, fallen.leaves = TRUE,
  main = "Decision Tree Visualization",
  box.palette = "RdBu", shadow.col = "gray", nn = TRUE)
```

Decision Tree Visualization



Checking if the cp value in the final model can be optimized and pruning the tree to avoid overfitting.

```
optimal_cp <- tree_model$bestTune$cp
pruned_tree <- prune(tree_model$finalModel, cp = optimal_cp)
```

Plotting the pruned tree if pruning improved performance

```
rpart.plot(pruned_tree, type = 2, extra = 101, fallen.leaves = TRUE,
  main = "Pruned Decision Tree",
  box.palette = "RdBu", shadow.col = "gray", nn = TRUE)
```

Pruned Decision Tree

